

An Introduction to Unit Testing

- **A 'unit' is a small piece of your code**
 - A small piece of functionality
- **A unit test verifies the correctness of that unit of code**
 - A purist might say that in a well-written unit test, only a single 'thing' should be able to fail
 - Generally accepted rule-of-thumb: a unit test should take less than a second to complete

Why Write Unit Tests?

- **Unit testing provides verification that your code is functioning correctly**
- **Much faster than testing your entire program each time you modify the code**
 - Fastest MapReduce job on a cluster will take many seconds
 - Even in pseudo-distributed mode
 - Even running in LocalJobRunner mode will take several seconds
 - LocalJobRunner mode is discussed later in the course
 - Unit tests help you iterate faster in your code development

Why MRUnit?

- **JUnit is a popular Java unit testing framework**
- **Problem: JUnit cannot be used directly to test Mappers or Reducers**
 - Unit tests require mocking up classes in the MapReduce framework
 - A lot of work
- **MRUnit is built on top of JUnit**
 - Works with the mockito framework to provide required mock objects
- **Allows you to test your code from within an IDE**
 - Much easier to debug

JUnit Basics (1)

- **@Test**

- Java annotation
- Indicates that this method is a test which JUnit should execute

- **@Before**

- Java annotation
- Tells JUnit to call this method before every @Test method
 - Two @Test methods would result in the @Before method being called twice

JUnit Basics (2)

- **JUnit test methods:**

- `assertEquals()`, `assertNotNull()` etc
 - Fail if the conditions of the statement are not met
- `fail(msg)`
 - Explicitly fails the test with the given error message

- **With a JUnit test open in Eclipse, run all tests in the class by going to Run → Run**
- **Eclipse also provides functionality to run all JUnit tests in your project**
- **Other IDEs have similar functionality**

JUnit: Example Code

```
import static org.junit.Assert.assertEquals;

import org.junit.Before;
import org.junit.Test;

public class JUnitHelloWorld {
    protected String s;

    @Before
    public void setup() {
        s = "HELLO WORLD";
    }

    @Test
    public void testHelloWorldSuccess() {
        s = s.toLowerCase();
        assertEquals("hello world", s);
    }

    @Test
    public void testHelloWorldFail() {
        assertEquals("hello world", s);
    }
}
```

Using MRUnit to Test MapReduce Code

- MRUnit builds on top of JUnit
- Provides a mock `InputSplit` and other classes
- Can test just the Mapper, just the Reducer, or the full MapReduce flow

MRUnit: Example Code – Mapper Unit Test (1)

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.junit.Before;
import org.junit.Test;

public class TestWordCount {
    Mapper<LongWritable, Text, Text, IntWritable> mapper;

    @Before
    public void setUp() {
        WordMapper mapper = new WordMapper();
        mapper = new Mapper<LongWritable, Text, Text, IntWritable>();
        mapper.setMapper(mapper);
    }

    @Test
    public void testMapper() {
        mapper.withInput(new LongWritable(1), new Text("cat dog"));
        mapper.withOutput(new Text("cat"), new IntWritable(1));
        mapper.withOutput(new Text("dog"), new IntWritable(1));
        mapper.runTest();
    }
}
```


MRUnit: Example Code – Mapper Unit Test (2)

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.junit.Before;
import org.junit.Test;
```

```
public
```

```
Mapper
```

```
@Before
```

```
public
```

```
void
```

```
mapDriver = new MapDriver<LongWritable, Text, Text, IntWritable>();
mapDriver.setMapper(mapper);
```

```
}
```

```
@Test
```

```
public void testMapper() {
```

```
mapDriver.withInput(new LongWritable(1), new Text("cat dog"));
```

```
mapDriver.withOutput(new Text("cat"), new IntWritable(1));
```

```
mapDriver.withOutput(new Text("dog"), new IntWritable(1));
```

```
mapDriver.runTest();
```

```
}
```

```
}
```

Import the relevant JUnit classes and the MRUnit MapDriver class as we will be writing a unit test for our Mapper.

MRUnit: Example Code – Mapper Unit Test (3)

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.junit.Before;
import org.junit.Test;
```

```
public class TestWordCount {
    MapDriver<LongWritable, Text, Text, IntWritable> mapDriver;
```

MapDriver is an MRUnit class (not a user-defined driver).

```
@Before
public void setUp() {
    mapDriver = new MapDriver<LongWritable, Text, Text, IntWritable>();
    mapDriver.setMapper(new Mapper());
}

@Test
public void testMapper() {
    mapDriver.withInput(new LongWritable(1), new Text("cat dog"));
    mapDriver.withOutput(new Text("cat"), new IntWritable(1));
    mapDriver.withOutput(new Text("dog"), new IntWritable(1));
    mapDriver.runTest();
}
}
```

MRUnit: Example Code – Mapper Unit Test (4)

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.junit.Before;
import org.junit.Test;

public class TestWordCount {
    Mapper<LongWritable, Text, Text, IntWritable> mapper;

    @Before
    public void setUp() {
        WordMapper mapper = new WordMapper();
        mapper = new Mapper<LongWritable, Text, Text, IntWritable>();
        mapper.setMapper(mapper);
    }

    @Test
    public void test() {
        mapper.runTest();
    }
}
```

Set up the test. This method will be called before every test, just as with JUnit.

MRUnit: Example Code – Mapper Unit Test (5)

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.junit.Before;
import org.junit.Test;

public class TestWordCount {
    Mapper<LongWritable, Text, Text, IntWritable> mapDriver;

    @Before
    public void setUp() {
        mapDriver = new WordCountMapper();
    }

    @Test
    public void testMapper() {
        mapDriver.withInput(new LongWritable(1), new Text("cat dog"));
        mapDriver.withOutput(new Text("cat"), new IntWritable(1));
        mapDriver.withOutput(new Text("dog"), new IntWritable(1));
        mapDriver.runTest();
    }
}
```

The test itself. Note that the order in which the output is specified is important – it must match the order in which the output will be created by the Mapper.

MRUnit Drivers (1)

- **MRUnit has a MapDriver, a ReduceDriver, and a MapReduceDriver**
- **Methods to specify test input and output:**
 - withInput
 - Specifies input to the Mapper/Reducer
 - Builder method that can be chained
 - withOutput
 - Specifies expected output from the Mapper/Reducer
 - Builder method that can be chained
 - addOutput
 - Similar to withOutput but returns void

MRUnit Drivers (2)

- **Methods to run tests:**

- `runTest`

- Runs the test and verifies the output

- `run`

- Runs the test and returns the result set

- Ignores previous `withOutput` and `addOutput` calls

- **Drivers take a single (key, value) pair as input**

- **Can take multiple (key, value) pairs as expected output**

- **If you are calling `driver.runTest()` or `driver.run()` multiple times, call `driver.resetOutput()` between each call**

- MRUnit will fail if you do not do this

Running Unit Tests From Eclipse

