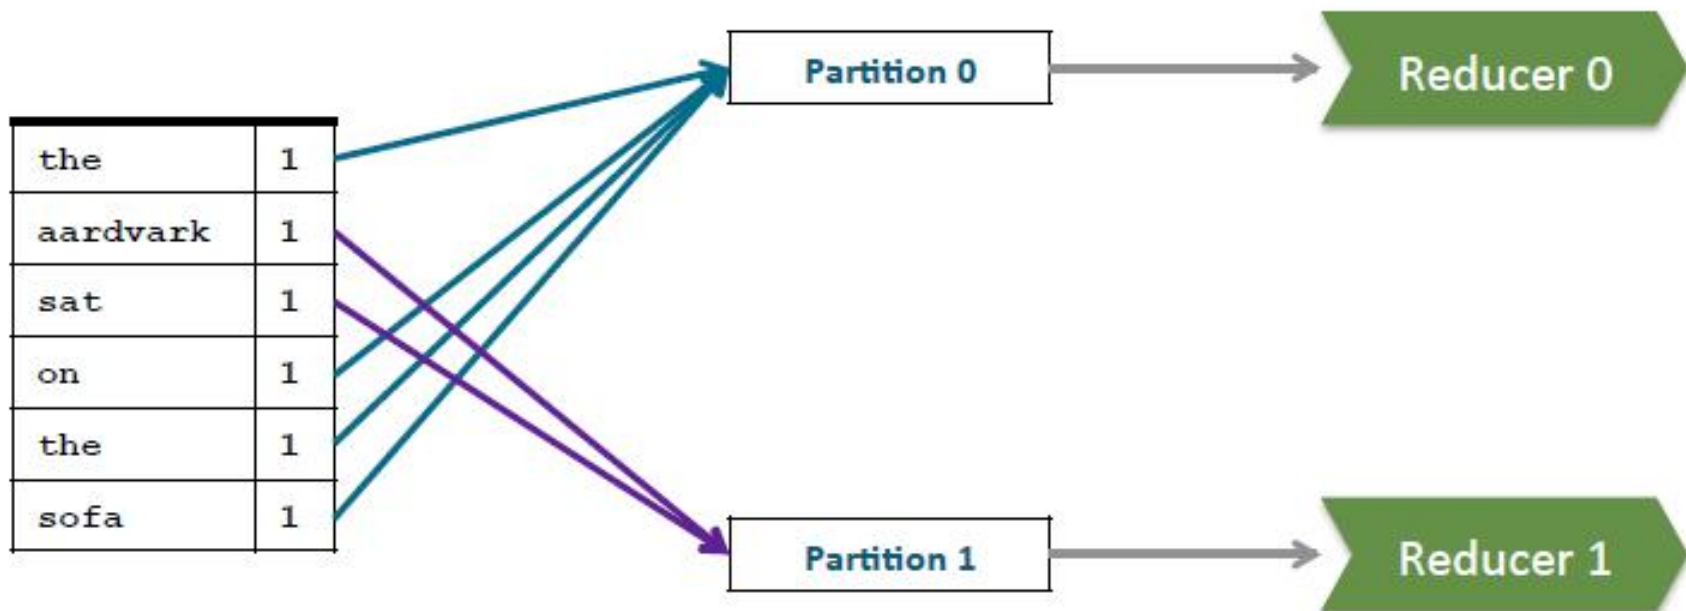


What Does the Partitioner Do?

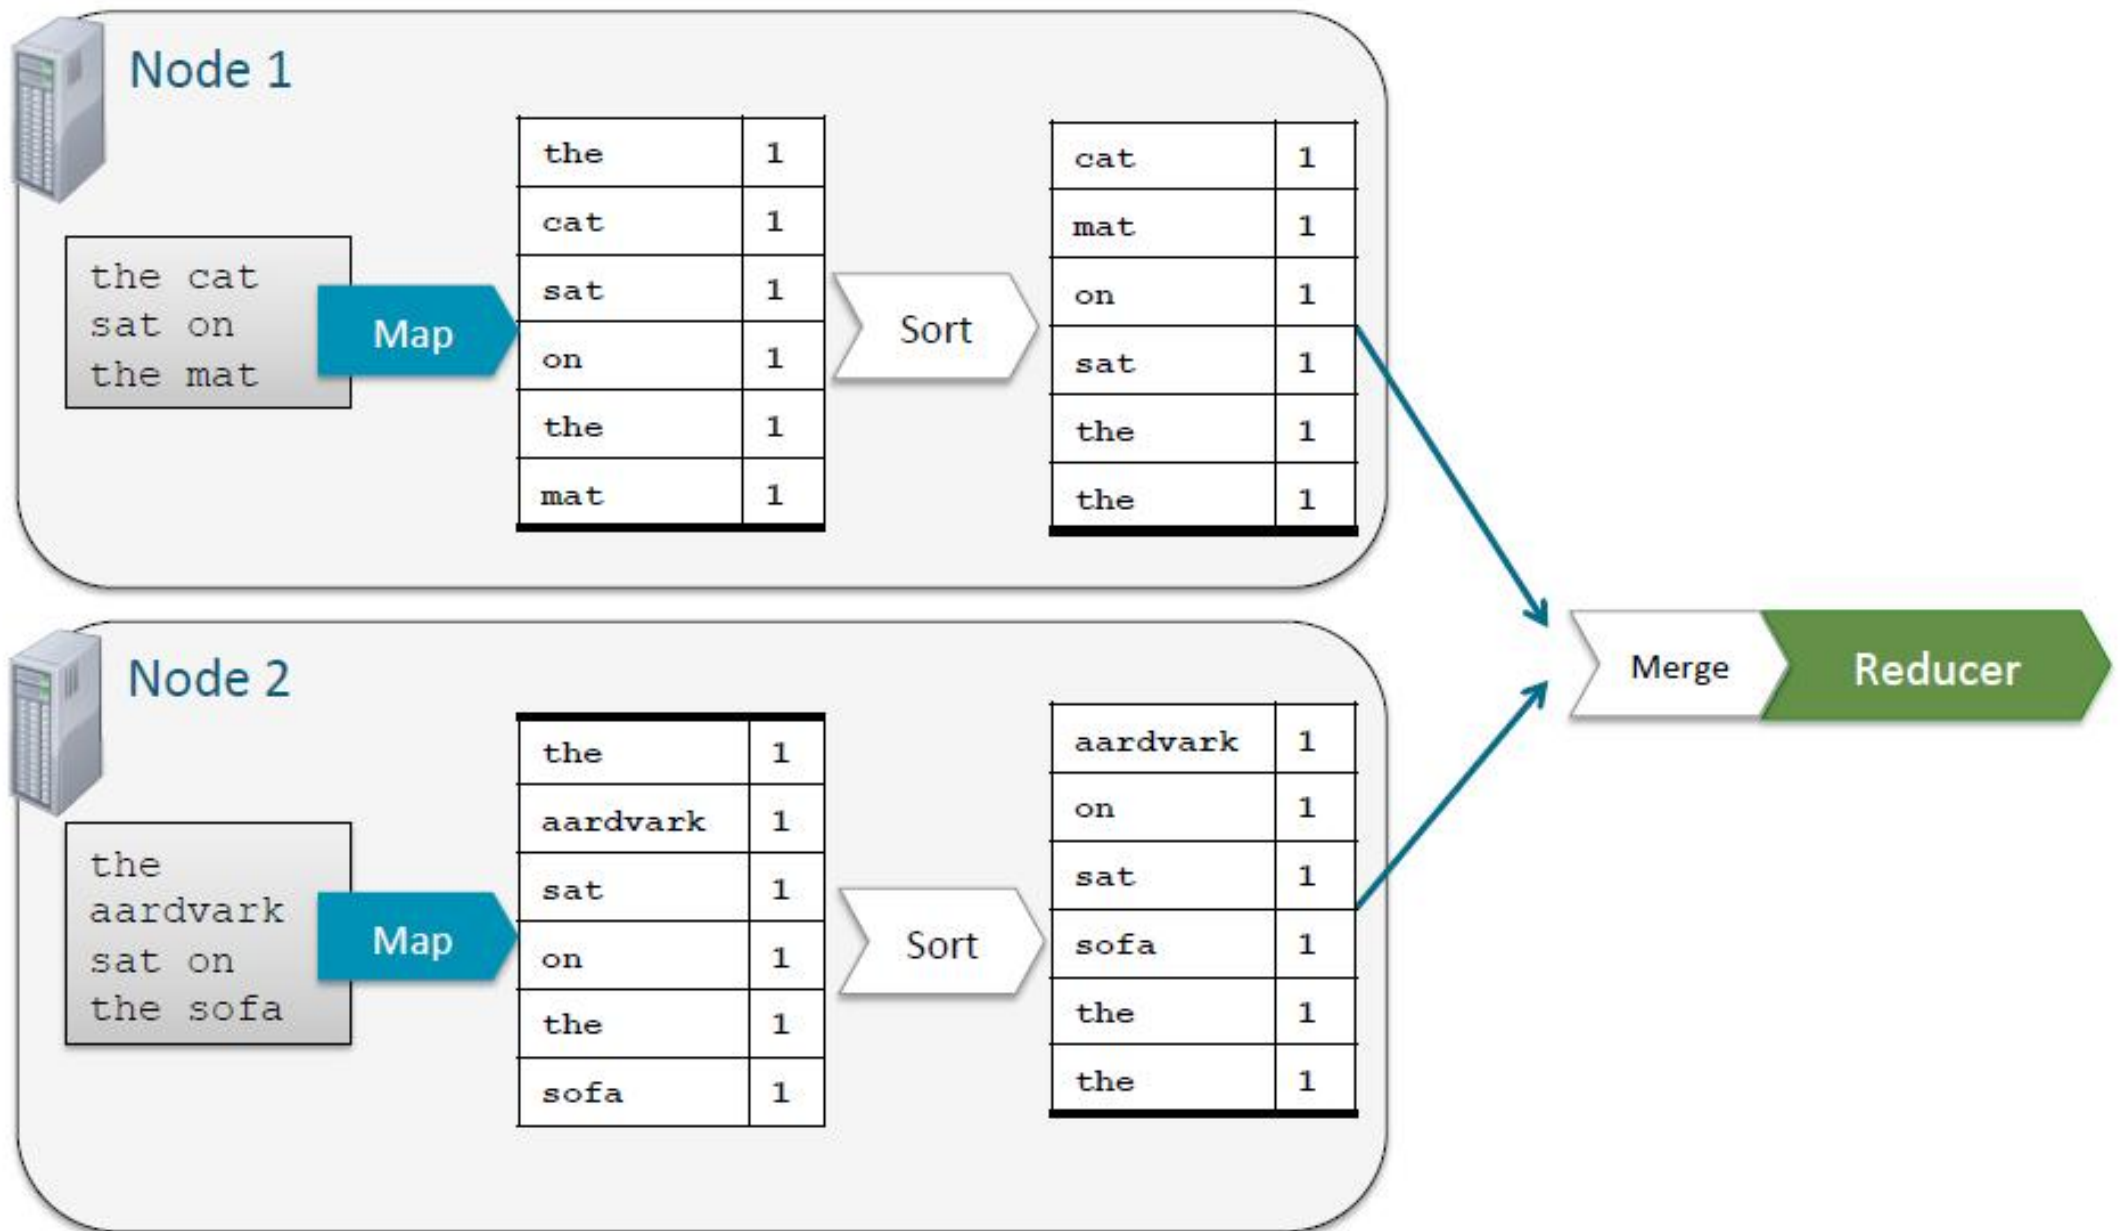
- The Partitioner determines which Reducer each intermediate key and its associated values goes to

getPartition:

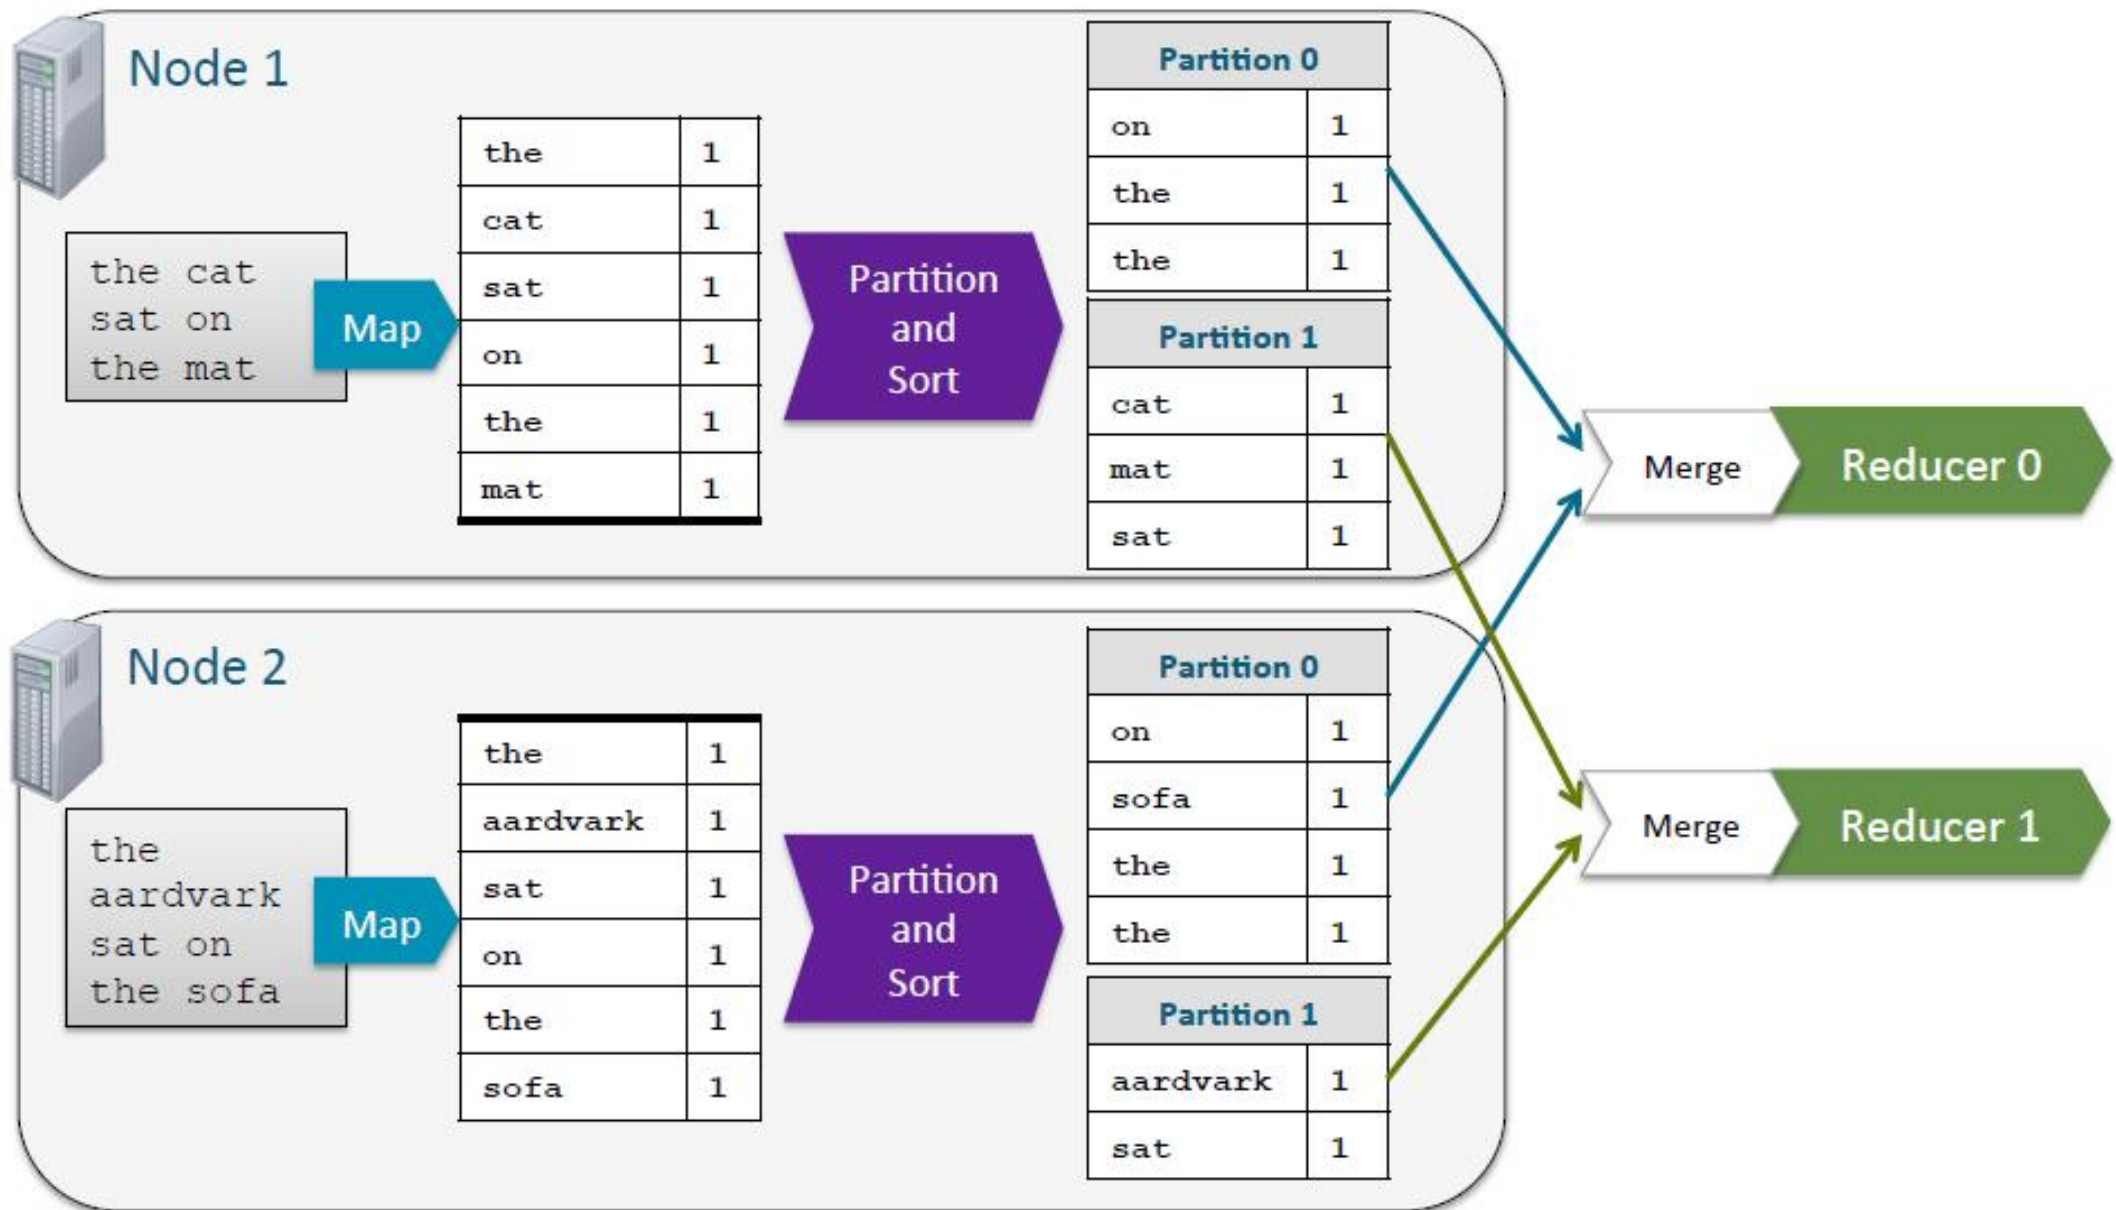
(inter_key, inter_value, num_reducers) → partition



Example: WordCount with One Reducer



Example: WordCount with Two Reducers



The Default Partitioner

- **The default Partitioner is the HashPartitioner**
 - Uses the Java `hashCode` method
 - Guarantees all pairs with the same key go to the same Reducer

```
public class HashPartitioner<K, V> extends Partitioner<K, V> {  
  
    public int getPartition(K key, V value, int numReduceTasks) {  
        return (key.hashCode() % numReduceTasks);  
    }  
}
```


How Many Reducers Do You Need?

- **An important consideration when creating your job is to determine the number of Reducers specified**
- **Default is a single Reducer**
- **With a single Reducer, one task receives *all* keys in sorted order**
 - This is sometimes advantageous if the output must be in completely sorted order
 - Can cause significant problems if there is a large amount of intermediate data
 - Node on which the Reducer is running may not have enough disk space to hold all intermediate data
 - The Reducer will take a long time to run

Jobs Which Require a Single Reducer

- **If a job needs to output a file where all keys are listed in sorted order, a single Reducer must be used**
- **Alternatively, the `TotalOrderPartitioner` can be used**
 - Uses an externally generated file which contains information about intermediate key distribution
 - Partitions data such that all keys which go to the first Reducer are smaller than any which go to the second, etc
 - In this way, multiple Reducers can be used
 - Concatenating the Reducers' output files results in a totally ordered list

Jobs Which Require a Fixed Number of Reducers

- **Some jobs will require a specific number of Reducers**
- **Example: a job must output one file per day of the week**
 - Key will be the weekday
 - Seven Reducers will be specified
 - A Partitioner will be written which sends one key to each Reducer

Jobs With a Variable Number of Reducers (1)

- **Many jobs can be run with a variable number of Reducers**
- **Developer must decide how many to specify**
 - Each Reducer should get a reasonable amount of intermediate data, but not too much
- **Typical way to determine how many Reducers to specify:**
 - Test the job with a relatively small test data set
 - Extrapolate to calculate the amount of intermediate data expected from the 'real' input data
 - Use that to calculate the number of Reducers which should be specified

Custom Partitioners

- Sometimes you will need to write your own Partitioner
- Example: your key is a custom `WritableComparable` which contains a pair of values `(a, b)`
 - You may decide that all keys with the same value for `a` need to go to the same Reducer
 - The default Partitioner is not sufficient in this case

Custom Partitioners

- **Custom Partitioners are needed when performing a secondary sort**
- **Custom Partitioners are also useful to avoid potential performance issues**
 - To avoid one Reducer having to deal with many very large lists of values
 - Example: in our word count job, we wouldn't want a single Reducer dealing with all the three- and four-letter words, while another only had to handle 10- and 11-letter words

Creating a Custom Partitioner

1. **Create a class that extends Partitioner**
2. **Override the `getPartition` method**
 - Return an int between 0 and one less than the number of Reducers
 - e.g., if there are 10 Reducers, return an int between 0 and 9

```
import org.apache.hadoop.mapreduce.Partitioner;

public class MyPartitioner<K,V> extends Partitioner<K,V> {

    @Override
    public int getPartition(K key, V value, int numReduceTasks) {
        //determine reducer number between 0 and numReduceTasks-1
        //...
        return reducer;
    }
}
```


Using a Custom Partitioner

- **Specify the custom Partitioner in your driver code**

```
job.setPartitionerClass(MyPartitioner.class);
```

Aside: Setting up Variables for your Partitioner

- If you need to set up variables for use in your partitioner, it should implement `Configurable`
- If a Hadoop object implements `Configurable`, its `setConf()` method will be called once, when it is instantiated
- You can therefore set up variables in the `setConf()` method which your `getPartition()` method will then be able to access

Aside: Setting up Variables for your Partitioner

```
class MyPartitioner extends Partitioner<K, V> implements Configurable {  
  
    private Configuration configuration;  
    // Define your own variables here  
  
    @Override  
    public void setConf(Configuration configuration) {  
        this.configuration = configuration;  
        // Set up your variables here  
    }  
  
    @Override  
    public Configuration getConf() {  
        return configuration;  
    }  
  
    public int getPartition(K key, V value, int numReduceTasks) {  
        // Use variables here  
    }  
  
}
```